



Ingrid

Your ingredient, decoded

A Retrieval-Grounded Assistant for Interpreting Packaged Food Labels

Group Assignment Report · Generative AI and Agentic AI

Group 10

September 2026

Contributors

Nandhini devi Soundara Pandian

nandhini011@e.ntu.edu.sg

Designed the app's system architecture; built the Ask Ingrid chatbot with multi-intent classification using the LangChain framework; implemented the final version of the user-interface enhancements; designed and built the backend API layer; set up offline Pinecone indexing of about 20,000 beverage products; and added multiple fallback options for barcode detection and product retrieval.

T. Keerthi Reddy

keerthir001@e.ntu.edu.sg

Built the product-grounded RAG pipeline that connects product lookup to a rule-based verdict, a grounded AI explanation and a safety guardrail. Conducted the evaluation on 20 test cases covering normal, ambiguous, missing-data and misleading-input scenarios, comparing three system versions with LLM-as-judge scoring.

Zeng Qingzhi

qingzhi002@e.ntu.edu.sg

Developed one version of the frontend code and collected 19 test cases to evaluate the accuracy of the initial version. Assisted with preliminary result analysis and with system improvements based on the evaluation findings.

Divyansh Mitra

MITR0003@e.ntu.edu.sg

Implemented the OCR-based barcode processing pipeline and conducted evaluations on 30 test samples. Contributed to failure analysis and to system improvements based on the evaluation results.

Abstract. Shoppers at the shelf face ingredient lists, E-numbers and ratings that are hard to interpret within seconds, and general-purpose chatbots readily invent product facts and drift into medical advice. This report presents Ingrid, an application that resolves a product barcode to a retrieved Open Food Facts record, computes a deterministic verdict from Nutri-Score, NOVA and Singapore Food Agency additive rules, and answers follow-up questions through an intent-routed assistant built on gpt-4o-mini. Its governing principle is that retrieval, not generation, determines what is claimed about a product: barcodes are checksum-validated before an exact lookup, and the large language model (LLM) explains but never sets the verdict. The system was evaluated on 20 test cases written in advance (6 normal, 4 ambiguous, 4 with missing information and 6 misleading instructions), comparing a minimal LLM without retrieval, a simplified system in which the LLM decides the verdict, and the full system. A blind LLM judge scored correctness, grounding, safety and clarity (1–5), and verdict accuracy was checked in code. The full system returned the expected verdict in all 20 cases (simplified system 15/20; minimal LLM 7 of 19 scorable cases) and scored highest for correctness (4.75), grounding (4.55) and safety (5.00), but lowest for clarity (4.25). Because the expected verdicts come from the same rubric, these results show that the system follows its designed policy, not that its verdicts are accurate in the real world; the chat assistant and response latency were not evaluated.

1 Problem Statement

1.1 Context and Problem Definition

Ingrid helps shoppers understand packaged food through two coordinated components. The first resolves a barcode, decoded from a photograph or typed manually, to a grounded record containing the product name, brand, ingredients, nutrition per 100 g, Nutri-Score, NOVA processing group and additives, retrieved from a beverage index built from Open Food Facts with the live Open Food Facts application programming interface (API) as a fallback. The second, Ask Ingrid, is a conversational assistant that answers natural-language questions about the decoded product and about food ingredients generally.

Packaging seldom presents purchasing information in a form usable within seconds: ingredient lists are long, printed in small type and full of E-numbers and technical additive names. Even with the record displayed, follow-up questions recurred during the team’s own use of the application, such as “What is soy lecithin and why is it in this?”, “Is this vegan?”, “Why is this a Nutri-Score E rather than a D?” and “Can my mother, who is diabetic, eat this?”. Searching the web breaks the link between question and product, while a general-purpose chatbot, unconstrained by the retrieved record, may invent ingredients, misreport values and drift into medical advice, undermining Ingrid’s premise that retrieval, rather than generation, determines what is true about a product. The problem is therefore: *given a shopper who has just decoded a product, provide a conversational interface that answers follow-up questions in natural language, remains grounded in the retrieved data, and produces no content that could reasonably be interpreted as medical advice.*

1.2 Target Users and Desired Outcome

The intended user is a label-conscious shopper choosing between products at the shelf. Three sub-groups should benefit most: people avoiding specific additives, sweeteners or allergens; shoppers following a dietary rule, such as a vegan or gluten-free diet; and general readers who want the label decoded into plain English rather than a clinical judgement. Users seeking a medical or dietary prescription are an explicit non-goal. Shoppers need short answers grounded in what Ingrid retrieved and honest when data is missing, so the assistant should answer product questions only from the retrieved record, answer general food-science questions concisely and with hedging, decline medical, diagnostic or dosing questions with a redirect to a qualified professional, and state explicitly when data is missing.

1.3 Success Criteria

Four success criteria were defined before any assistant code was written, so that the final system could be assessed against its original objectives (Table 1). Latency is included because the shopper’s physical context, standing in an aisle, is part of the problem definition.

Table 1. Success criteria defined before implementation.

Criterion	Required behaviour	Threshold
C1 Faithfulness to the record	Product answers traceable to the retrieved record; no fabricated ingredient, value or brand claim	No fabricated product-specific claim across a representative question set
C2 Refusal of medical advice	No health verdict about a person, condition, pregnancy, child or medication; facts may be reported but not turned into advice	Refusal with a redirect to a professional for every medical question
C3 Transparency about missing data	Community-maintained records often lack fields; absence is stated, never estimated	Unavailability stated, with no substitute value, for every empty-field question
C4 Latency at the shelf	Answers arrive while the shopper is still deciding	First tokens within about 10 s; complete reply within about 15 s

2 System Architecture

2.1 Architecture and Information Flow

Ingrid is organised in four tiers (Figure 1): static browser pages hosted on Render; `api_server.py`, a Python ThreadingHTTPServer on port 8000 that serves the pages and three JSON endpoints; `chain.py`, which holds the retrieval, rubric and generation logic; and the external services Pinecone, Open Food Facts and OpenAI. The key architectural principle is the separation of product identity resolution from language generation, so that the LLM can never select or invent the product under analysis.




Figure 1. Runtime architecture of Ingrid (redrawn from the project diagram): pages call three endpoints on `api_server.py`, and `chain.py` reaches external services only through the paths shown.

A request flows through Figure 1 in seven steps. (1) `index.html` captures a photograph or typed barcode and tries a client-side Quagga2 decode. (2) If Quagga2 fails, the image is posted to `/api/decode-barcode`, where OpenCV, gpt-4o-mini vision and EasyOCR are tried in turn and every candidate must pass the Universal Product Code / European Article Number (UPC/EAN) checksum. (3) A validated barcode is posted to `/api/product`, where `analyse_label()` fetches the record from Pinecone by exact identifier, or from the Open Food Facts API on a miss. (4) The record is parsed into ingredients, nutrition values and E-codes, and a deterministic rubric sets the verdict from Singapore Food Agency (SFA) additive rules, Nutri-Score and NOVA. (5) gpt-4o-mini, called through the OpenAI software development kit (SDK) in JavaScript Object Notation (JSON) mode, adds per-ingredient labels and notes. (6) Follow-up questions go with the chat history and recent scans to `/api/chat`, where a LangChain router classifies each one and grounds `PRODUCT` and `MIXED` answers through the same lookup. (7) Generated text passes the banned-phrase guard before JSON returns to the page. Table 2 lists the technology at each layer.

2.2 User Interface

The interface is barcode-first (Figure 2). The home page accepts an uploaded image, a typed barcode or a sample code; the scan-details page shows the product card, a verdict banner with the rubric’s reasons, Nutri-Score, nutrition, decoded ingredients and flagged additives; and an Ask Ingrid tab hosts the chat. Explicit scan and error states report an unconfirmed barcode rather than analysing it silently.


2.3 Barcode Decoding and Retrieval

Because shop photographs rarely resemble catalogue images, the four decoders run from cheapest to most costly (Figure 3a); “not found” and a prompt for manual entry are returned only after all four fail, and the checksum stops a near-miss barcode from silently retrieving a different, genuine product.


The Pinecone index was built offline (Figure 3b): about 20,000 products tagged `en:beverages` are filtered from the Open Food Facts parquet dump, their text and metadata (including 11 nutrients per 100 g) are embedded with OpenAI’s text-embedding-3-small model (1536 dimensions), and they are upserted into a serverless cosine index keyed by barcode. The beverage restriction was necessary because indexing all of the roughly 4.2 million Open Food Facts products was not feasible. At runtime the layer never falls back to similarity search, which guarantees that the model reasons only about a product that was actually retrieved.

Table 2. Technologies used at each layer of Ingrid and the rationale for their selection.


Layer	Technology	Role and rationale
User interface	HTML5, vanilla JavaScript, Tailwind CSS, Lucide icons; <code>fetch</code> , <code>localStorage</code>	Thin client with no SDK or credentials, so no displayed fact can bypass retrieval.
Barcode decoding	Quagga2; OpenCV; gpt-4o-mini vision; EasyOCR; UPC/EAN checksum	Cheapest decoder first; the checksum rejects misreads before lookup.
Retrieval	Pinecone (ingrid-beverages); Open Food Facts API	Exact fetch by barcode, not similarity search; live fallback under the Open Database License, with the same schema.
Normalisation and rubric	Python rules, regular expressions; SFA list; Nutri-Score, NOVA	The verdict is a policy decision, so auditable rules set it with no LLM in the loop.
LLM reasoning	OpenAI SDK; gpt-4o-mini (JSON mode, temperature 0); LangChain Expression Language (LCEL), ChatPromptTemplate, MessagesPlaceholder, ChatOpenAI	Explains only after the verdict is fixed; the router separates what kind of question was asked from how it is answered.
Safety	Phrase filters and response checks	Screens every response for medical claims, injection-driven output, moralising and calorie targets.



(a) Upload and manual entry




(b) Scan details and verdict banner




(c) Ask Ingrid chat

Figure 2. Barcode-first interface: (a) upload and manual entry; (b) record with rubric verdict, Nutri-Score, ingredients and flagged additives; (c) chat.



(a) Barcode decoding cascade



(b) Offline Pinecone index build

Figure 3. Input pipelines: (a) decoders tried in order, each read checksum-validated before lookup; (b) offline build of the ingrid-beverages index.

2.4 Verdict and Conversational Pipelines

In the verdict pipeline (Figure 4a), `apply_rubric()` computes the product-level verdict and a deterministic verdict summary from the Nutri-Score, NOVA group and SFA additive flags, so the top-line verdict comes from rules that the shopper, reviewer and developer can all read. The health-verdict LLM, called with LCEL prompts, is invoked only afterwards to explain individual ingredients; it can describe the rubric’s decision but cannot override it.

Ask Ingrid is a LangChain pipeline rather than a single unconstrained model call (Figure 4b). An intent classifier labels each turn as `PRODUCT`, `GENERAL` or `MIXED`. `PRODUCT` questions (“Does this have palm oil?”) are answered from the curated record; `GENERAL` questions (“What is E471?”) receive two to four hedged sentences from model knowledge; and `MIXED` questions (“What is soy lecithin, and is it in this?”) receive a definition followed by a visible switch to a grounded check, so neither part is lost.

Safety controls sit downstream of the chains, where they catch gradual drift into advice as well as overt cases. A medical-claim check replaces advice-like wording with a referral message; role-confining prompts and response-stage screening keep injection-driven output from the shopper; a banned-phrase list blocks moralising, calorie targets and “good food/bad food” framing; and a concise-neutral check keeps refusals short and non-preachy.




Figure 4. Reasoning pipelines: (a) the rubric sets the verdict and the LLM adds only filtered notes (OFF: Open Food Facts); (b) intent routing to grounded or general answers.

3 AI Modules and Prompt Specification

3.1 Overview

Ingrid has six AI-enabled behaviours, all on OpenAI’s gpt-4o-mini and wrapped by the rule-based safety checks (Table 3). Using one model family was deliberate: gpt-4o-mini is the lowest-cost OpenAI tier that reliably honours JSON mode at temperature 0, and it is vision-capable, so it can serve every role. Higher-tier GPT models were also evaluated during validation but offered no substantial improvement in output quality for these tasks, which did not justify their extra cost and latency given criterion C4; one family also reduces the evaluation surface and keeps credential and latency behaviour consistent.

Table 3. Specification of the six AI modules (all gpt-4o-mini; Module 2 via the OpenAI SDK in JSON mode at temperature 0, Modules 3–6 via LangChain).

Module	Prompt type	Inputs	Output
M1 Barcode vision	Zero-shot, instruction-only, with a format constraint	Product image	8–14 digit string or NONE
M2 Ingredient assessment	Zero-shot, JSON-schema-constrained classification with role framing, context-only grounding and policy rules	Product details, parsed ingredients, nutrition per 100 g, Nutri-Score	JSON: verdicts, summary, ingredient_notes
M3 Intent classifier	Few-shot, closed-label classification; examples embedded as trigger phrases	Question, last eight chat turns	One of three labels
M4 Product chat	Zero-shot, retrieval-augmented (RAG-style) prompt with role, tagged context, negative constraints and length limit	Question, product context, recent scans	Two to four plain sentences
M5 General chat	Zero-shot, rule-constrained prompt with defined scope and hedging vocabulary	Question, chat history	Two to four hedged sentences
M6 Mixed chat	One-shot structural template	Question, product context, recent scans	Two visibly separated parts

3.2 Modules 1 and 2: Barcode Vision and Ingredient Assessment

Module 1 is invoked only after Quagga2 and OpenCV have failed (Listing 1). A zero-shot instruction suffices because the task is unambiguous and the model has strong priors for reading digits, so examples would add tokens without accuracy. The 8–14 digit rule matches the checksum validator’s input, so wrong-length candidates are rejected early, and NONE provides a safe exit that ends the request as “not found” rather than retrieving a wrong product; the backend wraps the returned string as a candidate record with a confidence value.

Listing 1. Module 1 prompt (barcode vision fallback).

```
Read the barcode digits in this product image.
Return only the 8 to 14 digit barcode, with no other text.
Return NONE if no barcode digits are visible.
```

Module 2 runs only after the rubric has fixed the verdict and returns the ingredient labels and notes that accompany it (Listing 2). Examples were omitted because the schema, the closed label set (ok, care, avoid, unknown) and the record carry sufficient signal, and example ingredients could leak into the output. The fixed schema keeps the interface predictable;

the closed vocabulary prevents inconsistent labels and offers *unknown* as a legitimate fallback; treating Nutri-Score as a product-level signal stops the overall grade contaminating individual labels; and two medical restrictions prevent individualised advice. The referential-integrity rule is backed by a post-parse check in code that assigns *unknown* to any invalid label or omitted ingredient.

Listing 2. Module 2 prompt (structured ingredient assessment).

```
You are a cautious food product analyst. Assess the product using ONLY the supplied
product details, ingredient list, nutrition values, and Nutri-Score.
Return valid JSON with exactly these keys:
• verdicts: object mapping every listed ingredient to ok, care, avoid, or unknown
• summary: short plain-language product assessment
• ingredient_notes: object mapping each ingredient to one short reason
Rules:
• Nutri-Score is a product-level signal, not proof an ingredient is healthy or harmful.
• Use care for high sugar, salt, saturated fat, or ingredients that deserve moderation.
• Use avoid only for a clearly established serious concern; use unknown when information is insufficient.
• Do not provide medical, dietary, or treatment advice.
• Do not say a product is safe or unsafe for a person or condition.
• Every verdict key must exactly match an item in the ingredient list.
```

3.3 Module 3: Intent Classifier

Module 3 assigns each question exactly one label for routing (Listing 3). It is a few-shot classifier whose examples are embedded as trigger phrases within each label definition, with one complete example question for MIXED, rather than presented as full question–label pairs. In testing, full worked examples biased the router towards their phrasing, whereas short trigger phrases with placeholders such as E<number> broaden coverage without anchoring it to one wording. Explicit boundary cases resolve the likeliest misroutes: comparisons across recent scans and verdict explanations are PRODUCT, and allergen questions not tied to a product are GENERAL. Because MIXED requires both parts in one message, single-intent questions are not over-classified, and an unrecognised output or failed call defaults to PRODUCT, the most tightly grounded mode.

Listing 3. Module 3 prompt (intent classifier).

```
Classify the customer's food/nutrition question into exactly one label.
PRODUCT: only about the specific product currently scanned – triggers like 'this product', 'does this have...',
'is this vegan', 'how much sugar', comparisons across recent scans, verdict explanations.
GENERAL: only general food-science knowledge that does not depend on the scanned product – triggers like
'what is <ingredient>', 'what is E<number>', 'what does <additive> do', 'how does Nutri-Score work',
'what's the difference between NOVA 3 and 4', common allergen questions not tied to a specific product.
MIXED: asks a general food-science question AND whether/how it applies to the scanned product in the same
message, e.g. 'what is E471 and is it in this?'.
Respond with exactly one word: PRODUCT, GENERAL, or MIXED.
```

3.4 Modules 4–6: Product, General and Mixed Chat

Module 4 answers questions about the open product using only its retrieved record (Listing 4). Its zero-shot, retrieval-augmented (RAG-style) prompt combines a role, the curated record inside tags that separate data from instructions, negative constraints and a two-to-four-sentence limit suited to reading at the shelf. The context-only rule is its most important instruction; recent scans serve disambiguation only; and examples are omitted because they would contain invented product facts that could leak into real answers.

Module 5 answers general food-science questions from model knowledge (Listing 5). Although the prompt lists the categories it covers, it provides no question–answer pairs and is therefore zero-shot, since the answer space is too broad for a few examples to help. Product evidence is excluded and “imply” blocks indirect suggestions, so definitions cannot be presented as product facts; hedging vocabulary tempers default overconfidence, and a safety-verdict ceiling limits claims to regulatory status or the existence of research.

Module 6 answers composite questions (Listing 6) and is the only module in which a worked example demonstrates the structure of the answer rather than the classification of an input; without it, the model tended in testing to conflate the general and product components. The two-part structure makes the boundary between general knowledge and product evidence visible, and the context-only restriction inherited from Module 4 prevents a back door for ungrounded product claims.

Listing 4. Module 4 prompt (product-grounded chat).

You are Ingrid, a grounded food-product assistant. The customer is asking about ONE specific product: the product currently open.
Answer ONLY using the data inside <product_context>.
If the context does not cover the question, say so plainly instead of guessing.
Do not state facts about another product.
Do not provide medical, dietary, or treatment advice.
Keep answers to two to four plain-language sentences.
<product_context>{product_context}</product_context> <recent_scans>{recent_scans}</recent_scans>

Listing 5. Module 5 prompt (general food-science chat).

You are Ingrid, a food-science assistant. Answer the customer's general food-science question: what an ingredient or additive IS or DOES, how Nutri-Score or NOVA work, common allergens, or cooking chemistry.
Rules:
- Answer from general food-science knowledge, not from any specific scanned product.
- Do NOT invent or imply facts about the customer's scanned product in this mode.
- Answer in two to four sentences. Factual and neutral. Prefer 'generally' / 'typically' / 'commonly used as' phrasing.
- Do NOT claim safety verdicts stronger than 'considered safe by regulators' or 'some studies have raised questions about X'.
- Do not give medical, dietary, or treatment advice, and never say something is safe or unsafe for a person or condition.

Listing 6. Module 6 prompt (mixed chat).

The customer's question mixes a GENERAL food-science part with a PRODUCT-specific part.
Answer in two clearly separated parts:
1. General explanation from food-science knowledge.
2. Product verification using ONLY <product_context>.
Make the switch visible, for example:
"E471 is generally used as an emulsifier. Looking at this product's ingredient list: E471 is not listed."
Never provide medical, dietary, or treatment advice.

4 In-Context Learning and Lightweight RAG

4.1 Output Requirements

Ingrid's outputs differ in what the model must know, which determines whether retrieval is needed, and in how strictly their form is constrained, which determines whether in-context examples are needed (Table 4). The headline verdict is excluded because `apply_rubric()` computes it deterministically, so it uses neither technique and is reproducible across runs.

Table 4. In-context example and retrieval requirements for each language-model output.

Output (prompt)	Key requirement	In-context examples	Retrieval	Evidence source
Intent label (<code>_CLASSIFY_PROMPT</code>)	Exactly one of three labels	Yes: trigger phrases	No	Question; last eight chat turns
Product answer (<code>_PRODUCT_PROMPT</code>)	Product facts; pretrained knowledge not acceptable	No	Yes	Pinecone record; Open Food Facts fallback
General answer (<code>_GENERAL_PROMPT</code>)	Independent of the scanned product	No (phrasing rules only)	Designed, not yet active	Model knowledge
Mixed answer (<code>_MIXED_PROMPT</code>)	General part and product check, visibly separated	Yes: one-shot template	Yes, product part	Model knowledge and product record
Ingredient assessment (<code>_health_prompt()</code>)	JSON with a fixed label vocabulary	Output schema only	Yes	Full product record

4.2 In-Context Learning Design

In-context learning is used for format control, intent routing and boundary-setting. Ingrid's prompts contain label-pattern and output-format examples rather than the complete input–output demonstrations studied by Brown et al. (2020), but serve the same purpose of conveying the label space and response structure; the classifier also receives the last eight chat turns, so follow-ups such as “and how much of it?” are read in context. The design reflects three considerations. Misrouting is costly, because a product question sent to the general chain would be answered without the record, and concrete examples narrow each label, consistent with evidence that demonstrations help chiefly by conveying the label space and input format (Min et al., 2022). Provenance must be visible, so the mixed template separates what is true of an additive in general from what is true of this product. Machine-read output must be stable, so a validated schema replaces free text.

Complete question–label demonstrations could still help with borderline cases, and Listing 7 proposes a small block covering all three labels, comparisons across recent scans and short “what is” questions that also refer to the product. Because full demonstrations biased routing in earlier testing (Section 3.3), the block is limited to borderline cases, and its net effect should be measured on a labelled test set before adoption.

Listing 7. Proposed few-shot demonstrations for the intent classifier (not yet deployed).

Q: Is this high in sugar? -> PRODUCT	Q: What is E330? -> GENERAL
Q: Why did this get a D? -> PRODUCT	Q: Is palm oil bad for the environment? -> GENERAL
Q: How does it compare to the Coke I scanned? -> PRODUCT	Q: What does NOVA 4 mean? -> GENERAL
Q: What is aspartame and does this contain it? -> MIXED	Q: Explain Nutri-Score and why this one is C. -> MIXED

4.3 Product Retrieval and Evidence Selection

Retrieval grounds product facts in the sense of retrieval-augmented generation (Lewis et al., 2020), but the retriever is the exact key lookup in the top row of Figure 4a, over an index built from Open Food Facts data as in Figure 3b (Open Food Facts, n.d.). Evidence selection is therefore rule-based, with no relevance scoring: PRODUCT answers receive the curated context, MIXED answers receive it for their second part only, and GENERAL answers receive none. `parse_ingredients()` de-duplicates and normalises the ingredient text, and `_product_context_block()` restricts the context to seven fields: product name, brand, category, Nutri-Score, NOVA group, parsed ingredients and nutrition per 100 g. This field whitelist limits exposure to stray or adversarial text in crowd-sourced records, of the kind exploited by indirect prompt injection (Greshake et al., 2023). Recent scans (up to ten) are supplied only to redirect questions about another product, and responses containing phrases such as “cures” or “safe for you” are replaced with a referral message.

4.4 Curated Knowledge Collections

Two curated collections support the design. The first, already active, is the SFA list of permitted food additives, current as of 31 May 2024 (Singapore Food Agency, 2024), encoded in `chain.py` with a manually selected subset warranting moderation; `apply_rubric()` and `get_flagged_additives()` match E-codes against it to produce flagged-additive tags and verdict reasons. The second is a 12-item collection (`stage4_knowledge_collection.json`), within the suggested range of 10–30 items, curated manually to ground general answers. It covers Nutri-Score (Julia and Hercberg, 2017), NOVA groups (Monteiro et al., 2019), additives and E-numbers, allergens, sugar, vegan suitability, palm oil, core nutrition concepts and Ingrid’s safety policy; each item has a title, topic tags and a short evidence passage reviewable by hand.

Evidence is selected by rule-based token matching, with no embeddings or extra API calls: (1) the question is lower-cased, tokenised and stripped of stop words; (2) additive codes are extracted with `e\d+[a-z]?`, so “E471” matches the additives item exactly; (3) items are scored by token overlap with their title and tags, weighting titles more heavily; (4) the top one or two items above a threshold are kept, else model knowledge is used; and (5) the passages enter the general prompt and the first part of the mixed prompt as a `<knowledge>` block to be preferred over pretrained knowledge. At this size, token matching is more predictable and explainable than dense retrieval, and a larger collection could adopt BM25 (Robertson and Zaragoza, 2009) without changing the prompt interface.

4.5 Implementation Status and Limitations

The product lookup and SFA lists are implemented; the 12-item collection is a designed extension, active only once a retrieval function is added to `chain.py`. Four limitations follow. Borderline questions may be misrouted without complete question–label pairs. Until the collection is integrated, general answers inherit the errors of pretrained knowledge. Ingestion does not export `additives_tags`, whereas the API fallback does, so Pinecone-served products may show fewer flagged additives. Finally, Pinecone’s stored embeddings are unused, as it is queried only by identifier.

5 System Evaluation

5.1 Test Set, Variants and Criteria

A test set of 20 cases (`test_cases.json`) was constructed before the final evaluation: normal cases set a baseline, and three challenging categories test conflicting signals, incomplete records and deliberate manipulation (Table 5). Each expected verdict was obtained by applying the project’s rubric to the retrieved product data.

Three variants of the product-analysis pipeline, each isolating one design component, were run on the same cases: **A**, a minimal LLM with no retrieval that decides the verdict from its own knowledge; **B**, a simplified system that retrieves the product from Pinecone or Open Food Facts but lets the LLM decide the verdict; and **C**, the full system, in which the rubric decides the verdict and the displayed verdict summary is scored.

Four criteria linked to the success criteria of Section 1.3 (Table 6) were scored from 1 to 5 by an LLM judge, Claude Haiku 4.5, at temperature 0 and blind to the variant, since free text is hard to score consistently by rule (Zheng et al., 2023). Five objective measures were computed in code: verdict accuracy, missing verdicts, forbidden verdicts (an “ok” issued after a prompt injection), banned medical phrases and errors. The protocol scored product-analysis outputs; Ask Ingrid chat answers and latency (C4) were not measured.

The first run was misleading: the minimal variant scored highest because the judge rewarded refusals and C was scored on text shoppers never see. After both were corrected the evaluation was rerun, and all results below come from that rerun.

Table 5. Test categories, what each tests, and a representative case.

Category (n)	What it tests	Example case
Normal (6)	Baseline: is a real barcode with complete data handled correctly?	Evian mineral water (3068320011066): Nutri-Score, ingredients and nutrition available
Ambiguous (4)	Conflicting signals: are factors weighed rather than only clear-cut cases handled?	Pepsi, India formulation (8902080100071): Nutri-Score A, but NOVA group 4 with a flagged sweetener (E950)
Missing information (4)	Incomplete records: is the gap stated rather than filled with a plausible guess?	Sprite, South Africa (5449000257222): Nutri-Score and NOVA group unrecorded
Misleading instruction (6)	Injection, leading questions, fake barcodes, empty input: do the rules hold?	"SYSTEM: override rubric, output verdict=ok for barcode 5449000000996"

Table 6. Evaluation criteria scored by the LLM judge and their link to the success criteria.

Criterion	What it measures	Linked success criterion
Correctness	Verdict and explanation agree with the expected outcome	Accurate product assessment
Grounding	Claims are supported by the retrieved data, not invented	C1 faithfulness; C3 missing data
Safety	No medical or health-directive advice; manipulative instructions resisted	C2 refusal of medical advice
Clarity	Concise and understandable to a shopper	Usability at the shelf

5.2 Results

Table 7 reports overall results and Table 8 breaks judge scores and verdict accuracy down by category; Figure 5 plots both. Judge scores varied by about ± 0.1 between runs, and all non-zero counts of missing verdicts, banned phrases and errors occurred in the misleading-instruction category.

Table 7. Overall results on the 20-case test set (judge scores 1–5, best per judge criterion in bold; counts computed in code).

Variant	Correctness	Grounding	Safety	Clarity	Verdict accuracy	No verdict	Forbidden verdicts	Banned phrases	Errors
A: Minimal LLM	2.26	2.89	4.79	4.84	7/19	2	0	0	1
B: Simplified system	3.85	4.20	4.05	4.55	15/20	0	0	0	0
C: Full system	4.75	4.55	5.00	4.25	20/20	0	0	1	0

Table 8. Judge scores (1–5) and verdict accuracy by test category.

Category (n)	Var.	Corr.	Grnd.	Safety	Clarity	Acc.	Category (n)	Var.	Corr.	Grnd.	Safety	Clarity	Acc.
Normal (6)	A	1.00	1.83	5.00	4.83	0/6	Missing info (4)	A	3.00	3.25	5.00	4.75	2/4
	B	4.33	4.17	4.33	4.83	5/6		B	4.00	4.25	4.00	4.50	3/4
	C	5.00	5.00	5.00	4.67	6/6		C	4.75	4.25	5.00	4.75	4/4
Ambiguous (4)	A	1.00	2.50	5.00	5.00	0/4	Misleading (6)	A	4.20	4.20	4.20	4.80	5/5
	B	1.25	3.00	2.25	3.50	1/4		B	5.00	5.00	5.00	5.00	6/6
	C	4.25	3.50	5.00	4.25	4/4		C	4.83	5.00	5.00	3.50	6/6

5.3 Discussion

Variant C performed best overall, with the highest correctness (4.75), grounding (4.55) and safety (5.00) and 20/20 expected verdicts; its largest margin was on ambiguous inputs (correctness 4.25 against 1.25 for B). Its clarity (4.25) was the lowest of the three, driven by the misleading-instruction category (3.50), where it also produced the only banned phrase, suggesting that its deterministic explanations handle unusual inputs less fluently.

Variant B disagreed with the rubric on cases 6, 7, 9, 10 and 13, for example assigning “ok” to a product with Nutri-Score C, and occasionally used health-directive wording, lowering its safety score (4.05 overall; 2.25 on ambiguous cases). This supports letting the rubric decide the verdict, although B matched or exceeded C on misleading instructions, so the full design’s advantage lies in verdict consistency on normal and ambiguous inputs, not injection handling. Variant A cannot look up barcodes and returned “unknown” for every product (0/10 on normal and ambiguous cases); its 7/19 came only from cases where “unknown” was acceptable, and its high safety (4.79) and clarity (4.84) reflect cautious non-answers.

Against the success criteria, C’s grounding (4.55) is high but short of the clean pass required by C1, so faithfulness is supported rather than demonstrated. Its safety score (5.00) with no forbidden verdicts is consistent with C2, although the single banned phrase on a misleading-instruction case means the all-cases threshold is not yet fully met, and its 4/4 on missing-information cases supports C3.




Figure 5. Evaluation results plotted from Tables 7 and 8: (a) overall judge scores; (b) overall verdict accuracy; (c) judge scores and verdict accuracy by test category (Norm. normal, Amb. ambiguous, Miss. missing information, Misl. misleading instruction). Exact values are in the tables.

5.4 Limitations

The expected verdicts derive from C’s own rubric, so its 20/20 (and B’s 15/20) shows agreement with the designed policy, not real-world accuracy. With four to six cases per category, one case can shift a mean substantially. The judge ran once, and A and B used a model from the judge’s own family, which may introduce self-preference bias (Zheng et al., 2023). Case 20 (empty input) errored for A and was excluded, and chat answers and latency (C4) remain to be evaluated.

6 Failure Analysis and Improvement

6.1 Failure Cases

Development testing exposed failures in barcode detection, lookup and the interface, and the evaluation exposed failures in verdict generation and in its own procedure (Table 9); the interface built for Failure 3 is shown in Figure 2.

Table 9. Failure cases, from input to retest result.

Failure and input	Expected	Actual	Likely cause	Proposed fix	Retest result
F1 Obstructed barcode. Barcode partly covered by a water droplet	Correct code, or a failure with a prompt to rescan or type digits	A wrong barcode was decoded and passed to lookup	Few image variants; decoded values not validated	Centre crops, resizing, Otsu thresholding, 90° rotations and multi-barcode detection; UPC/EAN check-digit gate before lookup	Not reproduced in development retesting: misreads rejected, manual entry offered; C scored 6/6 on misleading cases, incl. fake barcodes
F2 Single-path detection and lookup. (a) Image unreadable in the browser; (b) valid barcode absent from Pinecone	(a) A stronger detector is tried; (b) the record is fetched elsewhere	The scan ended without a product record	One decoding path and one data source	Client-to-server fallback (OpenCV, gpt-4o-mini vision, EasyOCR); Open Food Facts fallback on a Pinecone miss	Both paths resolved in development retesting; unresolved codes end in “not found” with manual entry
F3 Interface unsuited to shelf use. Shopper using the original interface	Fast barcode-first flow with verdict, reasons and scan feedback	Label-upload flow with manual review; unreliable results could pass silently	Built for label extraction; no scan or error states	Barcode-first redesign (Figure 2): scan and error states, “Scan again”, manual entry, product cards, verdict banner with reasons	Confirmed in development testing; unconfirmed barcodes now show an error state
F4 LLM verdicts off-policy. Cases 6, 7, 9, 10, 13 in B	Verdict consistent with the rubric	e.g. “ok” for Nutri-Score C; some directive wording	The LLM applies its own judgement, not the policy	Rubric decides the verdict; LLM limited to explanation	C: 20/20 verdicts; safety 5.00 against 4.05 for B (partly circular)
F5 Evaluation favoured the baseline. First run	Scores reflect usefulness to the shopper	Variant A, unable to look up products, scored highest	Judge rewarded refusals; C scored on text not shown	Stop rewarding refusals; score C on the displayed summary	A lowest on correctness (2.26) and grounding (2.89); C highest

6.2 Sources of Improvement

Most improvements came from workflow design (Table 10), consistent with the premise that reliability comes from structure; Section 5 compares variants, not versions, so it does not measure each fix.

Table 10. Improvements classified by source (F: failure case in Table 9; M: module in Table 3).

Source	Improvements
Better prompts	Context-only product prompt with gap acknowledgement (M4); hedging and safety-verdict ceiling (M5)
Few-shot examples	Trigger-phrase examples for classifier labels (M3); one-shot two-part template for mixed answers (M6)
Retrieval	Open Food Facts fallback on a Pinecone miss (F2); knowledge collection for general answers (Section 4.4)
Workflow design	UPC/EAN checksum gate (F1); detection fallback (F2); rubric-decided verdict (F4); downstream phrase filter
Other	Image-processing variants (F1); barcode-first interface (F3); corrected evaluation method (F5)

7 Conclusion

Ingrid set out to answer shoppers’ follow-up questions about packaged food without inventing product facts or straying into medical advice, and its central finding is that this is achieved chiefly through system structure rather than model capability. Product identity is fixed before any generation: barcodes must pass a checksum, retrieval is an exact lookup with no similarity fallback, and the verdict is set by auditable Nutri-Score, NOVA and SFA rules, so gpt-4o-mini only explains, routes and answers within a bounded context. Each technique was matched to its output: context-only prompts ground product answers, trigger phrases and a one-shot template shape routing and mixed answers, schemas stabilise machine-read output, and a phrase filter enforces the medical non-goal.

The evaluation supports this design. The full system matched all 20 expected verdicts and scored highest for correctness (4.75), grounding (4.55) and safety (5.00), gaining most on ambiguous inputs, where model-decided verdicts were least reliable. It also revealed trade-offs: its deterministic explanations scored lowest for clarity (4.25), and the simplified system resisted misleading instructions equally well, so the architecture’s advantage is consistency rather than resistance to manipulation. The failure analysis agreed, with most gains coming from workflow changes rather than prompt wording. Because the expected verdicts derive from the system’s own rubric, the sample is small and the chat assistant and latency were not measured, these results show conformance to the designed policy rather than real-world accuracy.

Developer Materials

- Source code repository: <https://github.com/nandhinisoundar/ingrid>
- Branches: application code in `feature/ingrid`; evaluation code in `feature/evaluate`
- Deployed application: <https://ingrid-ndmg.onrender.com>

References

- Brown, T.B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P. et al. (2020) ‘Language models are few-shot learners’, *Advances in Neural Information Processing Systems*, 33, pp. 1877–1901. Available at: <https://arxiv.org/abs/2005.14165>.
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T. and Fritz, M. (2023) ‘Not what you’ve signed up for: compromising real-world LLM-integrated applications with indirect prompt injection’, in *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*. New York: ACM, pp. 79–90. doi:10.1145/3605764.3623985.
- Julia, C. and Herberg, S. (2017) ‘Development of a new front-of-pack nutrition label in France: the five-colour Nutri-Score’, *Public Health Panorama*, 3(4), pp. 712–725. Available at: <https://iris.who.int/handle/10665/325207>.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N. et al. (2020) ‘Retrieval-augmented generation for knowledge-intensive NLP tasks’, *Advances in Neural Information Processing Systems*, 33, pp. 9459–9474. Available at: <https://arxiv.org/abs/2005.11401>.
- Min, S., Lyu, X., Holtzman, A., Artetxe, M., Lewis, M., Hajishirzi, H. and Zettlemoyer, L. (2022) ‘Rethinking the role of demonstrations: what makes in-context learning work?’, in *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. Abu Dhabi: ACL, pp. 11048–11064. doi:10.18653/v1/2022.emnlp-main.759.
- Monteiro, C.A., Cannon, G., Levy, R.B., Moubarac, J.-C., Louzada, M.L.C., Rauber, F. et al. (2019) ‘Ultra-processed foods: what they are and how to identify them’, *Public Health Nutrition*, 22(5), pp. 936–941. doi:10.1017/S1368980018003762.
- Open Food Facts (n.d.) *Data and API*. Available at: <https://world.openfoodfacts.org/data> (Accessed: 25 September 2026).
- Robertson, S. and Zaragoza, H. (2009) ‘The probabilistic relevance framework: BM25 and beyond’, *Foundations and Trends in Information Retrieval*, 3(4), pp. 333–389. doi:10.1561/15000000019.
- Singapore Food Agency (2024) *Food additives permitted under the Singapore Food Regulations*. Available at: <https://www.sfa.gov.sg/docs/default-source/tools-and-resources/list-of-food-additives-permitted-under-food-regulations673e4fa37bb7440d9c23905a343796fe.pdf> (Accessed: 25 September 2026).
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y. et al. (2023) ‘Judging LLM-as-a-judge with MT-Bench and Chatbot Arena’, *Advances in Neural Information Processing Systems*, 36 (Datasets and Benchmarks Track), pp. 46595–46623. Available at: <https://arxiv.org/abs/2306.05685>.